

Report

February 24, 2023

1 Assignment 1 Report

1.0.1 Approach and design choices

In section 2 we were asked to manually backpropagate the error from the loss function. We were given $\text{delta} = -2 * (y - \hat{y})$ in formula (5). The following steps, undo activation \rightarrow calculate weight & bias gradient \rightarrow calculate previous activated layer \rightarrow ::, can be calculated in a loop until we reach the input layer. The challenge with this section was to make sure the orientation of the matrices were aligned for the matrix operations.

In section 3 we were asked to create a ML pipeline where we manually updated the weights with weight decay and momentum. First I needed to understand how the dataset was built up. I looked at a sample and saw that it contained an image and an integer and the samples were equally distributed across the integers 0-9. I displayed an image from each integer to understand what each number meant and saw that I needed to use categories 0 and 2.

It was important to use a fixed seed when separating the training and validation sets. That way I could calculate the mean and std of only the training set to compose a transformer for preprocessing all the images. When preprocessing you should only use the training data to prevent information from the validation/testing data to bleed into your training set.

The documentation in `torch.optim.SGD` was helpful when implementing weight decay and momentum and I decided to use a dictionary with parameter names as keys to keep track of the momentum.

1.0.2 Questions

- a) In section 2 we are asked to manually propagate the loss. This can be done in PyTorch with the `.backwards()` function.
- b) The `autograd.gradcheck` in pytorch is a method used to check if the computed gradient of a function seems correct. I could use this method to double check my computed gradients for the activation functions.
- c) The `.step()` function of the optimizer in PyTorch normally handles updating the weights.
- d) The SGD algorithm without momentum can get stuck in local minima and take a long time to get there. Momentum is used to make use of the history of the decent to help make better choices of how far to make the next step. The more we move in one direction the higher the momentum. Momentum is a term taken from physics and its use there is analogous to its use in ML.

- e) The purpose of regularization in SGD to mitigate the influence of outliers. L2 regularization punishes extreme outliers harshly so that the algorithm will prefer a “smoother curve”.
- f) The parameters I chose to test were every combination of $\text{lr} = [0.01, 0.005]$, $\text{decays} = [0, 0.001]$, $\text{momentums} = [0, 0.99]$. I wanted to have the option of $\text{momentum} = 0$ and $\text{decay} = 0$ in my sample and I wanted to try and vary each parameter with atleast two options. This meant that I tested $2^3 = 8$ different sets of parameters. The best model was the model with $\text{lr}=0.005$, $\text{decay}=0.001$, and $\text{momentum}=0.99$. I used accuracy as a evaluation metric to select the best performing model. These were the scores on the different datasets: Train accuracy: 0.9126 Val accuracy: 0.8433 Test accuracy: 0.8405
- g) With the best model selected the accuracy on the test set is about the same as the accuracy on the validation set. This gives me confidence that the performace on unseen data should be similar to the score 0.84. The performance is better on the training set (0.91). This indicates that the model hasn't generalized sufficiently. I believe training for more epochs would help with the generalization. 10 epochs are not suffucuent but I capped the training there because of time constraints.